

Programming Assignment 2 – Basic Probability, Computing and Statistics 2015

Fall 2015, Master of Logic, University of Amsterdam

Submission deadline: Monday, September 14th, 9 a.m.

Note: if the assignment is unclear to you or if you get stuck, do not hesitate to contact [Philip](#).

1 Exercise

This week we will start to work with some real data. The website [Project Gutenberg](#) provides access to thousands of books which have run out of copyright. They are freely downloadable for you to read. Our goal is to analyse these books for some of their basic properties. In particular, we want to find their most frequent and least frequent words and also find the counts for specific words. For example, we might want to know how often the word *Alice* occurs in *Alice in Wonderland*.

Just as last week we will again write an interactive program. This time, our program will be able to read text files from the command line. Moreover, it will keep interacting with the user until the user explicitly tells it to stop. Notice that during development you can experiment with any book you like. However, when grading we will all use [this version of War and Peace](#).

As you will notice, the file contains some document information from Project Gutenberg. This means that our word counts will not be entirely correct as there are also words that do not belong to the main text. We will not be bothered by this, however.

One last remark: when you download any book of your choice, please make sure to download the **Plain Text UTF-8** version of it. Otherwise your program will not be able to read the file (at least not without some additional effort).

Here is what your program must be able to do:

- Read a text file from the command line upon start (the path to the text file will be provided as a command line argument or after the program reminds the user)

- Ask the user what he wants to do. The user should be able to choose an operation by typing the appropriate number:
 1. Look for the most frequent words
 2. Look for the least frequent words
 3. Look for the frequency of a specific word
 4. Exit
- After each operation has been finished, the program should show the user the initial choice again and be able to perform further operations until the user decides to exit. Notice that reading the text anew for every operation would be extremely wasteful. You should try to get your program to memorize all relevant information.

2 Grading

Your program counts as executable if it starts up without having a path to the text file specified on the command line. In that case the program should explicitly ask the user to supply one.

- 1 point The program asks the user for the correct path to the file if that path is not supplied as a command line argument.
- 1 point The program gracefully shuts down if the supplied path does not exist. This means that the program prints a message informing the user that it will exit because the file was not found. It does not raise an error!
- 2 points The program reads the text file only once and stores all relevant information before showing the user any options. This has to be checked analytically by inspecting the code. Depending on your computer this operation can take several seconds.
- 2 points Whenever the user enters an invalid value, the program responds “Sorry, I do not understand this. Please choose again” and lets the user re-enter the value.
- 2 points When the user asks for a list of most frequent words, the program responds by asking for the number n of most frequent words to produce. It then outputs the n most frequent words to the screen as a list in which each line looks as follows:

word wordFrequency

If there are several words of the same frequency, they are ordered alphabetically.

- 2 points When the user asks for a list of least frequent words, the program responds by asking for the number n of least frequent words to produce. It then outputs the n least frequent words to the screen as a list in which each line looks as follows:

word wordFrequency

If there are several words of the same frequency, they are ordered alphabetically.

- 1 point The frequency of specific words is printed in the same format. If a word does not occur in the text, the program prints *Sorry, this word does not occur in the text.*

- 3 points If the program crashes at any point while testing.

Notice: for the frequencies, lists of the 1000 most frequent and 1000 least frequent words will be provided to you during the review period. We will also give you detailed instructions on how to compare two lists. For the specific words, you can just pick a couple of words from these lists.