

Programming Assignment 5 – Basic Probability, Computing and Statistics 2015

Fall 2015, Master of Logic, University of Amsterdam

Submission deadline: Monday, October 5th, 9 a.m.

Note: if the assignment is unclear to you or if you get stuck, do not hesitate to contact [Philip](#).

1 Naïve Bayes for Text Classification

This week you will achieve two of the big goals of this course. You will implement your first machine learning algorithm and run it on a real data set. The algorithm will be naïve Bayes. You will have to implement parameter learning by maximum likelihood estimation and prediction.

Data set The data set is the 20 new groups data set. It consists of roughly 1000 post for each of 20 newsgroups. The posts are of varying length but usually do not exceed the size of an e-mail. The data set is a classic in text classification and much research had been done with it in the beginning and mid 2000s. Nowadays it is considered too small and not very representative (after all, you want to learn to classify text beyond what people talk about in the 20 newsgroups).

We have pre-processed the data for you. We removed all punctuation and also split it up into a training set, a development set and a test set. This split is standard in machine learning. The training set contains the largest bulk of data and you use it to train your algorithm. The development set is usually rather small and allows you to check your algorithms performance on unseen data while developing. The test set is the final criterion of success. In a research paper you report your results on the test set. It is meant to represent new, unseen data. This means that you should not run you algorithm on it till the very end.

You can find the training data [here](#) and the dev set [here](#). The test set will only be made available next week to simulate a situation in which you really do not know the data on which you are going to apply your algorithm.

Implementation When implementing naïve Bayes, we have to take care of two things:

1. Most documents in the test and development sets will contain words that our classifier has never seen before. Hence, the maximum likelihood estimate for such words will be $\theta_{word} = 0$. This means the entire document will have probability 0. This is obviously not what we want. Therefore we apply the following trick: after having counted the words in the entire training set, we remove the words that have only have a count of 1, add up their counts and assign them to the unknown word `__UNK__`. In the prediction step, whenever we encounter a word that we have not seen before, we just pretend it is `__UNK__`. This way we can assign probabilities to all word.
2. The probabilities of each document will be rather low, so low in fact that they might get zeroed out. For this reason we will work with log-probs instead. Notice that this turns all multiplications into additions.

2 Grading

The grading is going to be based on the performance of your classifier rather than on your code alone. Please do not stress out about this. Try to work together and test your methods frequently (e.g. write your own unit tests).

- 8 points if the classifier achieves an accuracy of **TBD**% on the test set. You should use the dev set accuracy as an indicator for your classifier's performance. The classifier should achieve **TBD**% accuracy on the dev set.
- 2 point if you add a command line option that allows to modify the threshold on the word count that determines what words get mapped to the unknown word.
- 3 extra points if you manage to improve the performance of your classifier on the test set. This can be achieved by incorporating more features. For example, you may want to estimate the probability of upper case letters or include pairs of words (so-called bigrams) as features. Notice that if you include bigrams you also want to include the unknown bigram, that fulfills the same function as the unknown word.